

---

# CONVOLUTIONAL NEURAL NETWORK FOR SCENE RECOGNITION


---

VAISHALI SHIVAKUMAR (6738581)

AVERI RAY (6734926)

AISHWARYA JAYANTHRAJ (6788691)

*Assignment Report for Applied Machine Learning (EEEM068)*



*Department of Electronic Engineering*  
Faculty of Engineering and Physical  
Sciences University of Surrey  
Guildford, Surrey, GU2 7XH, UK

## *Table of Contents*

<i>Abstract .....</i>	<i>4</i>
<i>1. Introduction.....</i>	<i>4</i>
1.1. <i>Background.....</i>	<i>4</i>
1.2. <i>Objectives .....</i>	<i>4</i>
<i>2. Dataset .....</i>	<i>4</i>
2.1. <i>Description of Dataset .....</i>	<i>4</i>
2.2. <i>Statistics and Characteristics of the Dataset.....</i>	<i>4</i>
<i>3. Methodology.....</i>	<i>5</i>
3.1. <i>Overview of the model architecture (ResNet-34).....</i>	<i>5</i>
3.2. <i>Preprocessing steps and data augmentation techniques .....</i>	<i>5</i>
3.3. <i>Training procedure and hyperparameters .....</i>	<i>5</i>
<i>4. Fine-tuning Hyperparameters .....</i>	<i>6</i>
4.1. <i>Description of the hyperparameters tuned .....</i>	<i>6</i>
4.2. <i>Explanation of the motivation behind each hyperparameter choice.....</i>	<i>6</i>
4.3. <i>Impact of different hyperparameter values on training behavior and accuracy.....</i>	<i>7</i>
4.3.1.Experiment 1 .....	7
4.3.2.Experiment 2 .....	7
4.3.3.Experiment 3 .....	8
4.3.4.Experiment 4 .....	8
4.3.5.Experiment 5 .....	9
4.3.6.Experiment 6 .....	9
4.3.7.Experiment 7 .....	10
4.3.8.Experiment 8 .....	10
<i>5. Training Behavior.....</i>	<i>11</i>
5.1. <i>Analysis of the model's convergence behavior during training.....</i>	<i>11</i>
5.2. <i>Discussion of training loss and validation accuracy curves.....</i>	<i>11</i>
5.3. <i>Observations on the model's learning patterns and stability.....</i>	<i>17</i>
<i>6. Accuracy Measures .....</i>	<i>18</i>
6.1. <i>Calculation and interpretation of top-1 and top-5 accuracies.....</i>	<i>18</i>
6.2. <i>Comparison of accuracies across different hyperparameter settings.....</i>	<i>18</i>
6.3. <i>Discussion of any trends or insights revealed by the accuracy measures.....</i>	<i>18</i>

7. <i>Confusion Matrices</i> .....	19
7.1. <i>Presentation and interpretation of confusion matrices</i> .....	19
7.2. <i>Analysis of common confusion patterns and Misclassifications</i> .....	21
7.3. <i>Identification of challenging or similar classes based on the confusion matrices</i> .....	21
8. <i>CNN Visualizations</i> .....	22
8.1. <i>Visualization Techniques</i> .....	22
8.2. <i>Interpretation of Visualizations</i> .....	23
9. <i>Conclusion</i> .....	23
9.1. <i>Summary of Findings</i> .....	23
9.2. <i>Evaluation of the effectiveness of fine-tuning hyperparameters</i> .....	23
9.3. <i>Implications of the results for scene classification tasks</i> .....	24
<i>Citations</i> .....	25
<i>Appendix</i> .....	25

## ***Abstract***

This project focuses on training a Convolutional Neural Network (CNN) to recognize 40 different types of scenes using the "Places2 simp" dataset. The dataset consists of 40,000 color images, with each subfolder representing a different scene category. The ResNet-34 model is utilized for transfer learning, where the last layers are modified to fit the classification task.

The project involves implementing a dataset and data loader for the "Places2 simp" dataset. Since the images are initially sized at 128x128 pixels and the ResNet-34 model requires images of size 224x224 pixels, appropriate transformations are applied to resize the images. The dataset is randomly split into training and validation sets, following the typical 80% training and 20% validation ratio.

Data augmentation techniques are employed on the training set to enhance its diversity. Random translations, rotations within a range of -20 to 20 degrees, and scalings are applied to augment the images. A CNN model is then constructed by removing the last layers of ResNet-34 and adding new layers suitable for the scene classification problem. Different learning rates are experimented with to optimize the learning process.

The CNN model is trained, and hyperparameters such as batch size, number of epochs, and learning rate are fine-tuned to improve training behavior and accuracy measures. The training progress is visualized using a tensor board. The model's performance is evaluated on the validation set by calculating top-1 and top-5 classification accuracies, and a confusion matrix is generated to identify confusing cases.

To gain further insight into the model's performance, a subset of correctly and wrongly classified images from the validation set is randomly selected. The top-5 scores, along with the actual and predicted class names, are displayed to understand how CNN has learned to classify the scenes.

In addition to the provided dataset, a test set of scene images is created using photographs from Guildford, Surrey, or the London area. Accuracy rates (standard and top-5) and a confusion matrix are computed for the test set. Random test images are selected, and their top-5 scores and corresponding class names are displayed.

Different hyperparameter values are explored to observe their impact on training behavior and accuracy measures. The project aims to achieve a

minimum accuracy rate of 45% and a top-5 accuracy rate of 75% on the validation set.

Overall, this project demonstrates the implementation of a CNN using transfer learning and data augmentation techniques to classify scenes in the "Places2 simp" dataset. It highlights the importance of hyperparameter tuning and evaluation metrics for assessing the model's performance.

## ***1. Introduction***

### ***1.1. Background***

This project focuses on scene classification using Convolutional Neural Networks (CNNs). CNNs are a popular deep learning architecture for image analysis tasks, capable of learning hierarchical features from input images. Scene classification involves assigning a label or category to an image based on the scene or environment depicted in the image. It has numerous applications, including image retrieval, autonomous navigation, and content-based image analysis.

The dataset used in this project is the "Places2 simp" dataset, which is a simplified version of the larger "Places2" database. It consists of 40,000 color images, with each image belonging to one of 40 different scene categories. The dataset presents a challenging task due to the variability and complexity of scenes.

### ***1.2. Objectives***

1. Dataset Preparation: Load and preprocess the "Places2 simp" dataset, including resizing images to the required input size, splitting the dataset into training and validation sets, and ensuring a balanced representation of classes.
2. Transfer Learning: Utilize the ResNet-34 model, a widely used CNN architecture, as the base model for transfer learning. Modify the last layers of ResNet-34 to adapt it to the scene classification task.
3. Training and Hyperparameter Tuning: Train the modified CNN model on the training set with data augmentation techniques. Experiment with different hyperparameters such as batch size, learning rate, and number of epochs to optimize training behavior and maximize accuracy.
4. Evaluation: Evaluate the trained model on the validation set by calculating top-1 and top-5 classification accuracies. Generate a confusion matrix to visualize the model's performance and identify confusing cases.

5. Interpretability: Select a subset of correctly and wrongly classified images from the validation set and analyze the top-5 scores to gain insights into how the model has learned to classify scenes.

6. Test Set Evaluation: Create a separate test set consisting of scene images from Guildford, Surrey, or the London area. Evaluate the model's performance on this test set using accuracy rates (standard and top-5) and generate a confusion matrix.

7. Hyperparameter Exploration: Experiment with different hyperparameter values to observe their impact on training behavior and accuracy measures. Fine-tune the model to achieve the desired minimum accuracy rates on the validation set.

The main objective of this project is to develop an accurate and robust CNN model for scene classification. By leveraging transfer learning, data augmentation, and hyperparameter tuning, the aim is to achieve high classification accuracies on both the validation and test sets. Additionally, the project aims to gain insights into the model's performance and interpretability through the analysis of top-5 scores and confusion matrices.

## **2. Dataset**

### **2.1. Description of Dataset**

The dataset used in this project is the "Places2 simp" dataset, which consists of images representing 40 different categories of scenes or places. Each category, such as "airport terminal," "bar," "museum," and "playground," is represented by a subfolder in the dataset. The dataset is a simplified version of the larger "Places2" database.

The dataset is provided as a compressed zip file, "Places2 simp.zip," which can be obtained from SurreyLearn. The images in the dataset are color images with dimensions of 128x128 pixels. The folder structure of the dataset facilitates the labeling of images based on their corresponding classes.

### **2.2. Statics and Characteristics of the dataset**

The "Places2 simp" dataset consists of a total of 40,000 images, with each category containing 1,000 images. This distribution ensures an equal representation of each scene or place category in the dataset. The dataset aims to capture a wide range of scenes, including various indoor and outdoor environments.

The dataset exhibits characteristics such as diversity and complexity in terms of scene types, lighting conditions, viewpoints, and possible occlusions. It

provides a challenging task for scene classification due to the variations present in the images.

3. Understanding the statistical properties and characteristics of the dataset is crucial for training and evaluating the CNN model effectively. Analyzing the class distribution, image variations, and other dataset-specific factors helps in making informed decisions about data preprocessing, model architecture, and training strategies.

## **4. Methodology**

### **4.1. Overview of the model architecture (Resnet-34)**

The chosen model architecture for this project is ResNet-34, which is a deep convolutional neural network. ResNet-34 has shown excellent performance in image recognition tasks. It consists of 34 layers, including residual blocks that enable the network to learn more effectively and overcome the degradation problem. The ResNet-34 model is available within the torchvision package and can be loaded using the `'torchvision.models.resnet34()'` command.

### **4.2. Preprocessing steps and data augmentation techniques**

To prepare the dataset for training, several preprocessing steps and data augmentation techniques are employed. Firstly, since the ResNet-34 model expects images of size 224x224, the images from the "Places2 simp" dataset, which are initially 128x128 pixels, need to be resized to 224x224. Transformation functions are incorporated to resize the images accordingly.

Data augmentation is applied to the training set to increase the diversity of the data and improve model generalization. Random translations in the X and Y directions, random rotations within a certain range (e.g., -20 to 20 degrees), and random scaling are introduced as augmentation techniques. These techniques enhance the model's ability to handle variations in position, orientation, and scale within the images.

### **4.3. Training procedure and hyperparameters**

The dataset is split into a training set and a validation set using the common 80% for training and 20% for validation ratio. The model is created by removing the last layers of ResNet-34 and adding new layers suitable for the specific classification problem of recognizing 40 different scene categories. The learning rate is increased in the newly added layers to facilitate faster learning compared to the transferred layers from ResNet-34.

The CNN model is trained using the training set, and hyperparameters such as batch size, number of epochs, and learning rate are tuned to optimize performance. The training progress can be monitored and visualized using tools like Tensor Board.

During testing, the top-1 and top-5 classification accuracies are calculated, along with the confusion matrix, to evaluate the model's performance on the validation set. The confusion matrix is plotted in a way that helps understand the cases where the model gets confused between different classes.

A separate test set of scene images is created using photographs from Guildford or the wider Surrey or London area. The test set should include at least 5 to 10 images per category. The accuracy rates (standard and top-5 rates) and confusion matrix are estimated for this test set as well. Additionally, randomly chosen test images are examined to display the top-5 scores and corresponding class names.

To improve training behavior and accuracy measures, different values of hyperparameters can be explored. The impact of the chosen hyperparameters on the results should be observed and analyzed.

It is important to note that the target accuracy rate on the validation set is expected to be at least 45%, while the top-5 accuracy rate should be at least 75%. These metrics serve as benchmarks to evaluate the effectiveness of the trained CNN model. conditions.

## **5. Fine-tuning Hyperparameters**

### **5.1. Description of the hyperparameters tuned:**

The fine-tuning of hyperparameters is described, which involves modifying and optimizing various parameters that affect the training process and the performance of the CNN model. The specific hyperparameters that are tuned in the project are listed and explained.

### **4.2. Explanation of the motivation behind each hyperparameter choice:**

When choosing hyperparameters such as batch size, learning rate, epochs, weight decay, and optimizers (Adam, SGD, and AMSGrad), the following considerations were made:

#### **1. Batch Size:**

- A larger batch size can lead to faster training because more samples are processed in parallel, utilizing computational power efficiently.

- However, a larger batch size requires more memory, and if the batch size is too large, it may not fit into the available memory.

- The batch size should be chosen to strike a balance between efficient GPU utilization and memory constraints.

#### **2. Learning Rate:**

- The learning rate determines the step size at each iteration during the optimization process.

- A larger learning rate allows for larger updates in the model parameters, potentially leading to faster convergence.

- However, a very large learning rate may cause the optimization process to overshoot the optimal solution or even diverge.

- On the other hand, a small learning rate may result in slow convergence or getting stuck in a suboptimal solution.

- It is common to start with a higher learning rate and then decrease it gradually during training (learning rate scheduling) to allow for finer adjustments as the training progresses.

#### **3. Number of Epochs:**

- The number of epochs determines the number of times the entire training dataset is passed through the model.

- Training for too few epochs may result in an underfit model that hasn't fully captured the underlying patterns in the data.

- Training for too many epochs may lead to overfitting, where the model becomes too specialized to the training data and performs poorly on unseen data.

- The number of epochs should be chosen based on the convergence behavior of the training and validation loss/accuracy curves.

#### **4. Weight Decay:**

- Weight decay (also known as L2 regularization) is a regularization technique that adds a penalty term to the loss function to prevent overfitting.

- It encourages the model to learn smaller weights, preventing the model from relying too much on specific features.

- A higher weight decay value results in stronger regularization, while a lower value allows the model to fit the training data more closely.

- The weight decay value should be tuned through experimentation to find the right balance between preventing overfitting and preserving model performance.

#### 5. Optimizers (Adam, SGD, and AMSGrad):

- Optimizers determine how the model parameters are updated based on the computed gradients during backpropagation.

- Adam (Adaptive Moment Estimation) combines the advantages of RMSProp and Momentum methods. It adapts the learning rate for each parameter based on the past gradients and squared gradients.

- SGD (Stochastic Gradient Descent) updates the parameters by taking small steps in the direction of the negative gradient of the loss function.

- AMSGrad is a modification of Adam that ensures the learning rate is not increased too much for specific parameters, preventing overshooting.

- The choice of optimizer depends on the specific problem, and it is often determined empirically through experimentation.

- Adam is commonly used due to its adaptive learning rate and efficient convergence, but SGD and AMSGrad can also be effective in certain scenarios.

The selection of these hyperparameters should be based on the characteristics of the dataset, the complexity of the model, and the available computational resources. It is often recommended to perform hyperparameter tuning using techniques like grid search or random search to find the optimal combination of hyperparameters for the specific task discussed.

#### 4.3. *Impact of different hyperparameter values on training behavior and accuracy:*

The impact of different hyperparameter values on the training behavior and accuracy of the CNN model is analyzed. This section presents the results of experiments conducted with different hyperparameter settings and discusses how the changes influenced the convergence behavior, training progress, and final accuracy of the model.

##### 4.3.1. *Experiment 1:*

In Experiment 1, the ResNet-34 model was trained and evaluated using specific hyperparameters and

performance metrics. Here's a breakdown of the information provided:

Table 1: Hyperparameters for Experiment 1

optimizer	Adam
Learning rate	0.001
Feature extractor learning rate	0.0001
Weight Decay	0.0005
Momentum	0.5
No of Epochs	10
Train-batch-size	300
Test-batch-size	1024

These hyperparameters define the configuration of the training process. Adam is used as the optimizer, with a learning rate of 0.001 for the fully connected layer and a lower learning rate of 0.0001 for the rest of the model. Weight decay (L2 regularization) is set to 0.0005 to control overfitting, and a momentum value of 0.5 is used. The training is performed for 10 epochs, with a batch size of 300 during training and 1024 during testing.

Table 2: Performance of resnet34 in Experiment 1

Average Loss	0.5463
Top-1 Accuracy	76.66%
Top-5 Accuracy	95.76%
Top-10 Accuracy	95.76%
Top-20 Accuracy	95.76%

These performance metrics indicate the results achieved by the ResNet-34 model after training and testing. The average loss is 0.5463, representing the average value of the loss function during training. The top-1 accuracy is 76.66%, which means that the correct class label was predicted for 76.66% of the test samples. The top-5 accuracy, top-10 accuracy, and top-20 accuracy are all 95.76%, indicating that the correct label was included in the top-k predictions for a higher number of samples.

Overall, based on these results, the ResNet-34 model achieved a moderate top-1 accuracy of 76.66% and performed well in providing accurate predictions within the top-5, top-10, and top-20 predictions with an accuracy of 95.76%.

##### 4.3.2. *Experiment 2: Varying Batch Size and Learning Rate*

In Experiment 2, the ResNet-34 model was trained and evaluated with different batch sizes and learning rates. Here's an explanation of the provided data:

Table 3: Hyperparameters for Experiment 2

Optimizer	Adam
Learning rate	0.003
Feature extractor learning rate	0.0003
Weight Decay	0.0005
Momentum	0.5
No of Epochs	10
Train-batch-size	150
Test-batch-size	590

The hyperparameters define the configuration of the training process for Experiment 2. Adam optimizer is used with a learning rate of 0.003 for the fully connected layer and a lower learning rate of 0.0003 for the rest of the model. Weight decay (L2 regularization) is set to 0.0005, and a momentum value of 0.5 is used. The training is performed for 10 epochs, with a batch size of 150 during training and 590 during testing.

Table 4: Performance of resnet34 in Experiment 2

Average Loss	1.0413
Top-1 Accuracy	66.54%
Top-5 Accuracy	93.03%
Top-10 Accuracy	93.03%
Top-20 Accuracy	93.03%

These performance metrics represent the results obtained from Experiment 2. The average loss is 1.0413, which indicates the average value of the loss function during training. The top-1 accuracy is 66.54%, meaning that the correct class label was predicted for 66.54% of the test samples. The top-5 accuracy, top-10 accuracy, and top-20 accuracy are all 93.03%, indicating that the correct label was included in the top-k predictions for a higher number of samples.

Based on these results, it can be observed that with the specified hyperparameters in Experiment 2, the ResNet-34 model achieved a lower top-1 accuracy of 66.54% compared to Experiment 1. However, it still performed reasonably well in terms of top-5, top-10, and top-20 accuracies, reaching 93.03% in each case. These results suggest that the choice of batch size and learning rate can significantly impact the model's performance, and a smaller batch size with a higher learning rate may result in slightly lower accuracy but faster training.

#### 4.3.3. Experiment 3: Varying Epoch and Learning Rate

In Experiment 3, the ResNet-34 model was trained and evaluated with varying epochs and learning rates. Here's an explanation of the provided data:

Table 5: Hyperparameters for Experiment 3

optimizer	Adam
Learning rate	0.003
Feature extractor learning rate	0.0003
Weight Decay	0.0005
Momentum	0.5
No of Epochs	15
Train-batch-size	300
Test-batch-size	1024

The hyperparameters define the configuration of the training process for Experiment 3. Adam optimizer is used with a learning rate of 0.003 for the fully connected layer and a lower learning rate of 0.0003 for the rest of the model. Weight decay (L2 regularization) is set to 0.0005, and a momentum value of 0.5 is used. The training is performed for 15 epochs, with a batch size of 300 during training and 1024 during testing.

Table 6: Performance of resnet34 in Experiment 3

Average Loss	0.5828
Top-1 Accuracy	72.28%
Top-5 Accuracy	94.77%
Top-10 Accuracy	94.77%
Top-20 Accuracy	94.77%

These performance metrics represent the results obtained from Experiment 3. The average loss is 0.5828, which indicates the average value of the loss function during training. The top-1 accuracy is 72.28%, meaning that the correct class label was predicted for 72.28% of the test samples. The top-5 accuracy, top-10 accuracy, and top-20 accuracy are all 94.77%, indicating that the correct label was included in the top-k predictions for a higher number of samples.

Based on these results, it can be observed that with the specified hyperparameters in Experiment 3, the ResNet-34 model achieved a higher top-1 accuracy of 72.28% compared to Experiment 2. Additionally, the top-5, top-10, and top-20 accuracies remained consistently high at 94.77%. These results suggest that increasing the number of epochs can lead to improved accuracy, as the model has more opportunities to learn from the training data.

#### 4.3.4. Experiment 4: Varying Epochs and batch size.

In Experiment 4, the ResNet-34 model was trained and evaluated with varying epochs and batch sizes. Here's an explanation of the provided data:

Table 7: Hyperparameters for Experiment 4



Optimizer	Adam
Learning rate	0.001
Feature extractor learning rate	0.0001
Weight Decay	0.0005
Momentum	0.5
No of Epochs	15
Train-batch-size	150
Test-batch-size	590

The hyperparameters define the configuration of the training process for Experiment 4. Adam optimizer is used with a learning rate of 0.001 for the fully connected layer and a lower learning rate of 0.0001 for the rest of the model. Weight decay (L2 regularization) is set to 0.0005, and a momentum value of 0.5 is used. The training is performed for 15 epochs, with a batch size of 150 during training and 590 during testing.

Table 8: Performance of resnet34 in Experiment 4

Average Loss	0.3693
Top-1 Accuracy	81.41%
Top-5 Accuracy	96.95%
Top-10 Accuracy	96.95%
Top-20 Accuracy	96.95%

These performance metrics represent the results obtained from Experiment 4. The average loss is 0.3693, which indicates the average value of the loss function during training. The top-1 accuracy is 81.41%, meaning that the correct class label was predicted for 81.41% of the test samples. The top-5 accuracy, top-10 accuracy, and top-20 accuracy are all 96.95%, indicating that the correct label was included in the top-k predictions for a higher number of samples.

Based on these results, it can be observed that with the specified hyperparameters in Experiment 4, the ResNet-34 model achieved a higher top-1 accuracy of 81.41% compared to the previous experiments. Additionally, the top-5, top-10, and top-20 accuracies remained consistently high at 96.95%. These results suggest that increasing both the number of epochs and the batch size can lead to improved accuracy, as the model has more training iterations and can benefit from a larger batch of samples during each iteration.

#### 4.3.5. Experiment 5: Changing optimizer to SGD, Learning Rate

In Experiment 5, the ResNet-34 model was trained and evaluated using a different optimizer (SGD) and learning rate. Here's an explanation of the provided data:

Table 9: Hyperparameters for Experiment 5

optimizer	SGD
Learning rate	0.003
Feature extractor learning rate	0.0003
Weight Decay	0.0005
Momentum	0.5
No of Epochs	10
Train-batch-size	300
Test-batch-size	1024

For Experiment 6, the SGD optimizer is used with a learning rate of 0.003 for the fully connected layer and a lower learning rate of 0.0003 for the rest of the model. The weight decay is set to 0.0005, and the momentum is 0.5. The model is trained for 10 epochs, with a batch size of 300 during training and 1024 during testing.

Table 10: Performance of resnet34 in Experiment 5

Average Loss	1.3826
Top-1 Accuracy	60.04%
Top-5 Accuracy	89.65%
Top-10 Accuracy	89.65%
Top-20 Accuracy	89.65%

The performance metrics in Table 10 represent the results obtained from Experiment 6. The average loss is 1.3826, indicating the average value of the loss function during training. The top-1 accuracy is 60.04%, indicating that the correct class label was predicted for 60.04% of the test samples. The top-5, top-10, and top-20 accuracies are all 89.65%, indicating that the correct label was included in the top-k predictions for a higher percentage of samples.

Based on these results, it can be observed that using the SGD optimizer and the specified learning rates in Experiment 6 led to a decrease in accuracy compared to the previous experiments. The top-1 accuracy dropped to 60.04%, suggesting that the model struggled to make accurate predictions. The top-5, top-10, and top-20 accuracies also decreased to 89.65%. This indicates that the choice of optimizer and learning rate had a negative impact on the model's performance, resulting in lower accuracy.

#### 4.3.6. Experiment 6: Changing optimizer to amsgrad

In Experiment 6, the ResNet-34 model was trained and evaluated using the AMSGrad optimizer. Here's an explanation of the provided data:

Table 11: Hyperparameters for Experiment 6

Optimizer	amsgrad
-----------	---------

Learning rate	0.003
Feature extractor learning rate	0.0003
Weight Decay	0.0005
Momentum	0.5
No of Epochs	15
Train-batch-size	300
Test-batch-size	1024

For Experiment 6, the AMSGrad optimizer is used with a learning rate of 0.003 for the fully connected layer and a lower learning rate of 0.0003 for the rest of the model. The weight decay is set to 0.0005, and the momentum is 0.5. The model is trained for 15 epochs, with a batch size of 300 during training and 1024 during testing.

Table 12: Performance of resnet34 in Experiment 6

Average Loss	0.5865
Top-1 Accuracy	74.24%
Top-5 Accuracy	94.85%
Top-10 Accuracy	94.85%
Top-20 Accuracy	94.85%

The performance metrics in Table 10 represent the results obtained from Experiment 6. The average loss is 0.5865, indicating the average value of the loss function during training. The top-1 accuracy is 74.24%, indicating that the correct class label was predicted for 74.24% of the test samples. The top-5, top-10, and top-20 accuracies are all 94.85%, suggesting that the correct label was included in the top-k predictions for a higher percentage of samples.

Based on these results, it can be observed that using the AMSGrad optimizer in Experiment 6 led to improved performance compared to the previous experiment (Experiment 5 with SGD optimizer). The top-1 accuracy increased to 74.24%, indicating better prediction capabilities. The top-5, top-10, and top-20 accuracies also increased to 94.85%, demonstrating that the model performed well in capturing the correct labels within the top-k predictions.

Overall, the use of the AMSGrad optimizer in Experiment 6 resulted in better performance and higher accuracy compared to the SGD optimizer in the previous experiment.

#### 4.3.7. Experiment 7: Performance of Resnet50

In Experiment 7, the ResNet-50 model was trained and evaluated. Here's an explanation of the provided data:

Table 13: Hyperparameters for Experiment 7

Optimizer	Adam
Learning rate	0.001
Feature extractor learning rate	0.0001
Weight Decay	0.0005
Momentum	0.5
No of Epochs	15
Train-batch-size	150
Test-batch-size	590

For Experiment 7, the ResNet-50 model was trained using the Adam optimizer. The learning rate for the fully connected layer was set to 0.001, and the learning rate for the rest of the model (feature extractor) was 0.0001. The weight decay was set to 0.0005, and the momentum was 0.5. The model was trained for 15 epochs with a batch size of 150 during training and 590 during testing.

Table 14: Performance of resnet34 in Experiment 7

Average Loss	0.5912
Top-1 Accuracy	77.05%
Top-5 Accuracy	96.16%
Top-10 Accuracy	96.16%
Top-20 Accuracy	96.16%

The performance metrics in Table 14 represent the results obtained from Experiment 7. The average loss is 0.5912, indicating the average value of the loss function during training. The top-1 accuracy is 77.05%, meaning that the correct class label was predicted for 77.05% of the test samples. The top-5, top-10, and top-20 accuracies are all 96.16%, indicating that the correct label was included in the top-k predictions for a higher percentage of samples.

Based on these results, it can be observed that using the ResNet-50 model in Experiment 7 led to improved performance compared to the ResNet-34 model used in previous experiments. The top-1 accuracy increased to 77.05%, demonstrating better prediction capabilities. The top-5, top-10, and top-20 accuracies also increased to 96.16%, indicating that the model performed well in capturing the correct labels within the top-k predictions.

Overall, Experiment 7 with the ResNet-50 model achieved higher accuracy compared to previous experiments, suggesting that the deeper architecture of ResNet-50 allowed for better feature extraction and classification.

#### 4.3.8. Experiment 8: Performance of Alexnet

In Experiment 8, the AlexNet model was trained and evaluated. Here's an explanation of the provided data:

Table 15: Hyperparameters for Experiment 8

Optimizer	Adam
Learning rate	0.001
Feature extractor learning rate	0.0001
Weight Decay	0.0005
Momentum	0.5
No of Epochs	15
Train-batch-size	150
Test-batch-size	590

For Experiment 8, the AlexNet model was trained using the Adam optimizer. The learning rate for the fully connected layer was set to 0.001, and the learning rate for the rest of the model (feature extractor) was 0.0001. The weight decay was set to 0.0005, and the momentum was 0.5. The model was trained for 15 epochs with a batch size of 150 during training and 590 during testing.

Table 16: Performance of resnet34 in Experiment 8

Average Loss	1.2086
Top-1 Accuracy	73.16%
Top-5 Accuracy	93.14%
Top-10 Accuracy	93.14%
Top-20 Accuracy	93.14%

The performance metrics in Table 16 represent the results obtained from Experiment 8. The average loss is 1.2086, indicating the average value of the loss function during training. The top-1 accuracy is 73.16%, meaning that the correct class label was predicted for 73.16% of the test samples. The top-5, top-10, and top-20 accuracies are all 93.14%, indicating that the correct label was included in the top-k predictions for a higher percentage of samples.

Based on these results, it can be observed that the AlexNet model in Experiment 8 achieved moderate performance in terms of accuracy. The top-1 accuracy of 73.16% indicates that the model correctly predicted the class label for approximately 73% of the test samples. The top-5, top-10, and top-20 accuracies of 93.14% demonstrate that the model included the correct label within the top-k predictions for a higher percentage of samples.

Overall, the performance of the AlexNet model in Experiment 8 suggests that it was able to learn and classify the scene categories to a reasonable extent, although it achieved slightly lower accuracy compared to the ResNet models in the previous experiments.

## 6. Training Behavior

### 6.1. Analysis of the model's convergence behavior during training:

In the provided code, the convergence behavior of the CNN model is analyzed by monitoring the training loss and validation accuracy metrics. The training loss measures how well the model is fitting the training data, while the validation accuracy indicates the model's performance on unseen validation data. By observing these metrics over different epochs, insights can be gained into the model's learning dynamics and training performance.

### 6.2. Discussion of training loss and validation accuracy curves

The training loss and validation accuracy curves are plotted and analyzed to understand the model's progress during training. The training loss curve shows the trend of the loss value decreasing over epochs. A decreasing loss indicates that the model is improving its fit to the training data. On the other hand, an increasing or fluctuating loss suggests that the model is struggling to learn the underlying patterns in the data.

The validation accuracy curve depicts the accuracy achieved by the model on the validation set at each epoch. A rising validation accuracy indicates that the model is generalizing well and performing better on unseen data. Conversely, a stagnant or decreasing validation accuracy suggests that the model may be overfitting or failing to generalize beyond the training set.

By examining the training loss and validation accuracy curves, it is possible to assess the model's convergence behavior. Ideally, one would observe a decreasing training loss and an increasing validation accuracy, signifying successful convergence and generalization of the model.

#### 6.2.1. Experiment 1:

The training loss and validation accuracy curves provide insights into the training progress and model performance during Experiment 1. The hyperparameters used in this experiment are specified in Table 1, and the performance results are given in Table 2.


During the training process, the ResNet34 model was trained for 10 epochs with a batch size of 300. The Adam optimizer was used with a learning rate of 0.001 for the fully connected layer and 0.0001 for the rest of the model. A weight decay of 0.0005 and a momentum of 0.5 were applied to control overfitting.

The training loss curve represents the change in the loss function over the course of training. It indicates how well the model is fitting the training data. In Experiment 1, the average loss achieved was 0.5463, which suggests that the model's training converged reasonably well. The decreasing trend of the training loss curve indicates that the model improved its ability to minimize the discrepancy between predicted and actual labels as training progressed.



The validation accuracy curve represents the accuracy of the model on a separate validation dataset during each epoch. In Experiment 1, the ResNet34 model achieved a top-1 accuracy of 76.66%, indicating that it correctly predicted the class label for approximately 76.66% of the validation samples. The top-5, top-10, and top-20 accuracies were also high, all at 95.76%. This demonstrates that the model performed well in including the correct label within the top-k predictions, with high accuracy for multiple prediction options.


The validation accuracy curve shows an increasing trend over the epochs, indicating that the model's performance improved as it learned more from the training data. The consistent and high validation accuracies suggest that the model successfully generalized to unseen validation samples and performed well in classifying the scene categories.




Overall, the training loss and validation accuracy curves for Experiment 1 demonstrate that the ResNet34 model, trained with the specified hyperparameters, achieved good convergence in terms of the training loss and high accuracy in classifying the scene categories during validation. This indicates the effectiveness of the chosen hyperparameters and the suitability of the ResNet34 model for the given task of scene classification.

### 6.2.2. Experiment 2:

The The training loss and validation accuracy curves provide insights into the training progress and model performance during Experiment 2. The hyperparameters used in this experiment are specified in Table 3, and the performance results are given in Table 4.



During the training process, the ResNet34 model was trained for 10 epochs with a batch size of 150. The Adam optimizer was used with a learning rate of 0.003 for the fully connected layer and 0.0003 for the rest of the model. A weight decay of 0.0005 and a momentum of 0.5 were applied to control overfitting.



The training loss curve represents the change in the loss function over the course of training. It indicates how well the model is fitting the training data. In Experiment 2, the average loss achieved was 1.0413,

which suggests that the model's training converged, but the loss is higher compared to Experiment 1. The increasing or fluctuating trend of the training loss curve indicates that the model struggled to minimize the discrepancy between predicted and actual labels as training progressed.

The validation accuracy curve represents the accuracy of the model on a separate validation dataset during each epoch. In Experiment 2, the ResNet34 model achieved a top-1 accuracy of 66.54%, which is lower compared to Experiment 1. The top-5, top-10, and top-20 accuracies were also lower at 93.03%. This indicates that the model's performance in correctly predicting the class labels decreased compared to Experiment 1.

The validation accuracy curve shows a fluctuating trend over the epochs, indicating that the model's performance varied during training and did not consistently improve. The lower validation accuracies suggest that the model struggled to generalize well to unseen validation samples and had difficulty classifying the scene categories accurately.

Overall, the training loss and validation accuracy curves for Experiment 2 demonstrate that the ResNet34 model, trained with the specified hyperparameters, had difficulties in convergence and achieved lower accuracies compared to Experiment 1. This suggests that the chosen hyperparameters, specifically the combination of learning rate and batch size, may not have been optimal for the given task of scene classification. Further adjustments and optimization of hyperparameters may be needed to improve the model's performance. follows:

### 6.2.3. Experiment 3:

In Experiment 3, the training loss and validation accuracy curves were analyzed to understand the performance of the ResNet34 model. The hyperparameters used in this experiment are outlined in Table 5, and the corresponding performance results are provided in Table 6.


The training process consisted of training the ResNet34 model for 15 epochs with a batch size of 300. The Adam optimizer was utilized, with a learning rate of 0.003 for the fully connected layer and 0.0003 for the remaining parts of the model. A weight decay value of 0.0005 and a momentum value of 0.5 were employed to mitigate overfitting.

The training loss curve demonstrates the change in the loss function throughout the training process. It reflects how effectively the model is fitting the training data. In Experiment 3, the average loss achieved was 0.5828, which indicates that the model's training converged to a lower loss compared to

Experiment 2. The decreasing trend of the training loss curve suggests that the model gradually improved its ability to minimize the difference between predicted and actual labels as the training progressed.



The validation accuracy curve represents the accuracy of the model on a separate validation dataset at each epoch. In Experiment 3, the ResNet34 model attained a top-1 accuracy of 72.28%. Although this is an improvement compared to Experiment 2, it is still lower than the performance achieved in Experiment 1. The top-5, top-10, and top-20 accuracies were also higher at 94.77%. This suggests that the model's performance in accurately classifying scene categories improved, but there is still room for enhancement.



The validation accuracy curve exhibits an increasing trend, indicating that the model's performance improved steadily over the epochs. The higher validation accuracies imply that the model achieved better generalization to unseen validation samples and performed well in classifying the scene categories accurately.


In summary, the training loss and validation accuracy curves for Experiment 3 demonstrate that the ResNet34 model, trained with the specified hyperparameters, improved in terms of convergence and achieved higher accuracies compared to Experiment 2. This suggests that the adjustments

made to the learning rate, batch size, and number of epochs positively influenced the model's performance. However, the model's performance in Experiment 3 still falls short of the accuracy achieved in Experiment 1, indicating that further optimization may be necessary to enhance the model's capabilities for scene classification.

#### 6.2.4. Experiment 4:

In Experiment 4, the training loss and validation accuracy curves were analyzed to evaluate the performance of the ResNet34 model. The hyperparameters used in this experiment are outlined in Table 7, and the corresponding performance results are provided in Table 8.


The ResNet34 model was trained for 15 epochs with a batch size of 150 using the Adam optimizer. The learning rate for the fully connected layer was set to 0.001, and the feature extractor learning rate was set to 0.0001. A weight decay value of 0.0005 and a momentum value of 0.5 were applied to regularize the model and control overfitting.



The training loss curve depicts the change in the loss function over the course of training. In Experiment 4, the average loss achieved was 0.3693, indicating that the model effectively minimized the discrepancy between predicted and actual labels during training. The decreasing trend of the training loss curve signifies that the model progressively improved its fitting to the training data, resulting in a lower overall loss.

The validation accuracy curve illustrates the accuracy of the model on a separate validation dataset at each epoch. In Experiment 4, the ResNet34 model attained a top-1 accuracy of 81.41%, indicating that it correctly classified the scene category in the validation dataset for 81.41% of the samples. The top-5, top-10, and top-20 accuracies were also high at 96.95%. This suggests that the model's performance in accurately predicting scene categories significantly improved.

The validation accuracy curve exhibits an increasing trend, indicating that the model's performance steadily improved as the training progressed. The higher validation accuracies imply that the model achieved better generalization and successfully classified unseen validation samples with a high degree of accuracy.



In summary, the training loss and validation accuracy curves for Experiment 4 demonstrate that the ResNet34 model, trained with the specified hyperparameters, achieved a low training loss and high validation accuracies. The model effectively learned the features of the dataset and successfully classified scene categories. The excellent performance, with a top-1 accuracy of 81.41% and top-5 accuracy of 96.95%, indicates that the model's training converged well and produced accurate predictions.


#### 6.2.5. Experiment 5:

In Experiment 5, the training loss and validation accuracy curves were analyzed to evaluate the performance of the ResNet34 model with the SGD optimizer. The hyperparameters used in this experiment are outlined in Table 9, and the corresponding performance results are provided in Table 10.


The ResNet34 model was trained using the SGD optimizer with a learning rate of 0.003 for the fully connected layer and a feature extractor learning rate of 0.0003. A weight decay value of 0.0005 and a momentum value of 0.5 were applied for regularization. The model was trained for 10 epochs with a batch size of 300 during training and a batch size of 1024 during testing.

The training loss curve represents the change in the loss function during the training process. In Experiment 5, the average loss achieved was 1.3826, indicating that the model struggled to minimize the loss. The increasing trend of the training loss curve suggests that the model may have had difficulty converging and fitting the training data effectively.

The validation accuracy curve shows the accuracy of the model on a separate validation dataset at each epoch. In Experiment 5, the ResNet34 model achieved a top-1 accuracy of 60.04%, indicating that it correctly classified the scene category for 60.04% of the validation samples. The top-5, top-10, and top-20 accuracies were also relatively lower at 89.65%. These results indicate that the model's performance in accurately classifying the validation samples was not as strong as in other experiments.



The validation accuracy curve shows a fluctuating trend, suggesting that the model's performance varied during training and did not consistently improve over epochs. The lower validation accuracies indicate that the ResNet34 model may have struggled to capture important features and patterns in the dataset and may have faced challenges in generalizing effectively.




In summary, the training loss and validation accuracy curves for Experiment 5 indicate that the ResNet34 model, trained with the specified hyperparameters and the SGD optimizer, had difficulty converging and achieving high accuracy in classifying the scene categories. The model's performance, as indicated by the training loss and validation accuracy, was relatively lower compared to other experiments. Further analysis and adjustments to the hyperparameters and optimization strategy may be

necessary to improve the model's performance in this experiment.


#### 6.2.6. Experiment 6:

In Experiment 6, the training loss and validation accuracy curves were analyzed to evaluate the performance of the ResNet34 model with different optimizer and learning rate settings. The hyperparameters used in this experiment are outlined in Table 9, and the corresponding performance results are provided in Table 10.

For this experiment, the optimizer was changed to SGD, and the learning rate for the fully connected layer was set to 0.003, while the feature extractor learning rate was set to 0.0003. A weight decay value of 0.0005 and a momentum value of 0.5 were applied for regularization and to control overfitting. The model was trained for 10 epochs with a batch size of 300 and tested with a batch size of 1024.



The training loss curve shows the change in the loss function during the training process. In Experiment 6, the average loss achieved was 1.3826, indicating that the model had a higher loss compared to other experiments. This suggests that the model had some difficulty in fitting the training data and might have experienced slower convergence during training.






The validation accuracy curve represents the accuracy of the model on a separate validation dataset at each epoch. In Experiment 6, the ResNet34 model achieved a top-1 accuracy of 60.04%, indicating that it correctly classified the scene category for 60.04% of the validation samples. The top-5, top-10, and top-20 accuracies were also lower at 89.65%. These results suggest that the model's performance was relatively lower compared to other experiments.

The validation accuracy curve shows a fluctuating trend, indicating that the model's performance varied across different epochs. The lower validation accuracies imply that the model struggled to generalize well and accurately classify unseen validation samples.

In summary, the training loss and validation accuracy curves for Experiment 6 suggest that the ResNet34 model, trained with SGD optimizer and specific learning rate settings, faced challenges in achieving lower loss and higher accuracy compared to other experiments. The higher training loss and lower validation accuracies indicate that the model's performance was not as successful in this experiment. It is likely that the chosen optimizer and learning rate settings did not lead to optimal convergence and generalization for the given dataset.


#### 6.2.7. Experiment 7:

In Experiment 7, the training loss and validation accuracy curves were analyzed to evaluate the performance of the ResNet50 model. The hyperparameters used in this experiment are outlined in Table 13, and the corresponding performance results are provided in Table 14.



For this experiment, the ResNet50 model was trained using the Adam optimizer with a learning rate of 0.001 for the fully connected layer and a feature extractor learning rate of 0.0001. A weight decay value of 0.0005 and a momentum value of 0.5 were applied for regularization. The model was trained for 15 epochs

with a batch size of 150 during training and a batch size of 590 during testing.



The training loss curve shows the change in the loss function during the training process. In Experiment 7, the average loss achieved was 0.5912, indicating that the model was able to effectively minimize the loss and fit the training data well. The decreasing trend of the training loss curve suggests that the model converged during training and learned to capture the patterns in the dataset.

The validation accuracy curve represents the accuracy of the model on a separate validation dataset at each epoch. In Experiment 7, the ResNet50 model achieved a top-1 accuracy of 77.05%, indicating that it correctly classified the scene category for 77.05% of the validation samples. The top-5, top-10, and top-20 accuracies were also relatively high at 96.16%. These results indicate that the ResNet50 model performed well in accurately classifying the validation samples.

The validation accuracy curve shows an increasing trend, suggesting that the model improved its performance over the course of training and was able to generalize well to unseen validation samples. The relatively high validation accuracies indicate that the ResNet50 model successfully captured the complex features and patterns present in the dataset.


In summary, the training loss and validation accuracy curves for Experiment 7 demonstrate that the ResNet50 model, trained with the specified hyperparameters, achieved good convergence and high accuracy in classifying the scene categories. The model's performance suggests its ability to effectively learn and generalize complex features, resulting in accurate predictions on unseen data.

#### 6.2.8. Experiment 8:


In Experiment 8, the training loss and validation accuracy curves were analyzed to evaluate the performance of the AlexNet model. The hyperparameters used in this experiment are outlined



in Table 15, and the corresponding performance results are provided in Table 16.



For this experiment, the AlexNet model was trained using the Adam optimizer with a learning rate of 0.001 for the fully connected layer and a feature extractor learning rate of 0.0001. A weight decay value of 0.0005 and a momentum value of 0.5 were applied for regularization. The model was trained for 15 epochs with a batch size of 150 during training and a batch size of 590 during testing.



The training loss curve represents the change in the loss function during the training process. In Experiment 8, the average loss achieved was 1.2086, indicating that the model was able to minimize the loss, although it was relatively higher compared to other experiments. The decreasing trend of the training loss curve suggests that the model converged during training and learned to fit the training data.

The validation accuracy curve shows the accuracy of the model on a separate validation dataset at each epoch. In Experiment 8, the AlexNet model achieved a top-1 accuracy of 73.16%, indicating that it correctly classified the scene category for 73.16% of the validation samples. The top-5, top-10, and top-20 accuracies were also relatively high at 93.14%. These results indicate that the AlexNet model performed well in accurately classifying the validation samples.

The validation accuracy curve shows a fluctuating trend, suggesting that the model's performance varied during training and did not consistently improve over epochs. The relatively high validation accuracies indicate that the AlexNet model captured important features and patterns present in the dataset, although it may have struggled to generalize as effectively as other models.

In summary, the training loss and validation accuracy curves for Experiment 8 indicate that the AlexNet model, trained with the specified hyperparameters, achieved a reasonable level of convergence and accuracy in classifying the scene categories. However, compared to other experiments, the model's performance in terms of both training loss and validation accuracy was relatively lower. Further optimization and fine-tuning of the model may be necessary to improve its performance.

### 6.3. Observations on the model's learning patterns and stability

Based on the provided code and the results from experiments 1-8, several observations can be made regarding the model's learning patterns and stability.

Across all experiments, it can be observed that the model struggled to achieve high accuracies and consistently obtained relatively high average loss values. This suggests that the model encountered difficulties in learning and generalizing effectively on the given dataset.

When analyzing the training loss and validation accuracy curves, it can be seen that the loss values did not consistently decrease over epochs, indicating that the model had challenges in minimizing the loss and fitting the training data. The validation accuracies also varied across experiments, with some experiments achieving higher accuracies than others. This suggests that the model's performance was sensitive to changes in hyperparameters and training configurations.

In terms of stability, it can be observed that the model's performance was not consistent across different experiments. Varying hyperparameters, such as learning rate, batch size, and optimizer, resulted in different levels of accuracy and loss. This indicates that the model's learning patterns were not stable and that the choice of hyperparameters significantly affected its performance.

Overall, the results obtained from experiments 1-8 indicate that the model's learning patterns were not optimal, and its stability was limited. Further investigation and refinement of the model architecture, hyperparameters, or training strategies are necessary to improve its performance on the given dataset. It may be beneficial to explore additional

techniques such as fine-tuning, regularization, or ensemble learning to enhance the model's learning capacity and improve its stability.


7. Accuracy Measures

7.1. Calculation and interpretation of top-1 and top-5 accuracies

The top-1 and top-5 accuracies are calculated to evaluate the model's classification performance. The top-1 accuracy represents the percentage of correctly predicted classes for a given sample, while the top-5 accuracy considers whether the correct class is within the top 5 predicted classes. These accuracy measures provide a comprehensive assessment of the model's ability to recognize the correct scene category.

7.2. Comparison of accuracies across different hyperparameters settings

In experiments 1 to 8, various hyperparameter settings were investigated to understand their influence on the model's accuracy measures. By comparing the accuracies obtained under different hyperparameter configurations, we can assess the effectiveness of specific choices and identify optimal settings for achieving higher accuracies.



Across the experiments, it can be observed that different hyperparameter combinations resulted in varying levels of accuracy. Experiment 4, with the hyperparameters of Adam optimizer, a learning rate of 0.001, feature extractor learning rate of 0.0001, weight decay of 0.0005, momentum of 0.5, 15 epochs, a train batch size of 150, and a test batch size of 590, achieved the highest accuracies.

In particular, experiment 4 achieved a top-1 accuracy of 81.41%, top-5 accuracy of 96.95%, top-10 accuracy of 96.95%, and top-20 accuracy of 96.95%. This indicates that the chosen hyperparameter settings in

experiment 4 were effective in improving the model's ability to classify images accurately.

On the other hand, experiments 2 and 5, which involved different combinations of learning rates and optimizer choices, resulted in lower accuracies compared to experiment 4. This suggests that the learning rate and optimizer selection play significant roles in determining the model's performance. Experiment 6, which used a different optimizer (amsgrad), also showed lower accuracies compared to experiment 4, further highlighting the importance of the optimizer choice.

Overall, the comparison of accuracies across experiments 1 to 8 emphasizes the impact of hyperparameter settings on the model's performance. It highlights the importance of selecting appropriate values for learning rate, optimizer, weight decay, momentum, and batch sizes to achieve higher accuracies. The findings from these experiments can guide the selection of optimal hyperparameter configurations for improving the model's accuracy on the given dataset.

7.3. Discussion of any trends or insights revealed by the accuracy measures.

The comparison of accuracies across experiments 1 to 8 provides valuable insights into the trends and patterns in the model's performance. By examining the accuracy measures obtained under different hyperparameter configurations, we can gain insights into the strengths and limitations of the trained CNN model.

Experiment 4, with its specific hyperparameter settings, achieved the highest accuracies among all the experiments. This indicates that the chosen combination of the Adam optimizer, a learning rate of 0.001, feature extractor learning rate of 0.0001, weight decay of 0.0005, momentum of 0.5, 15 epochs, a train batch size of 150, and a test batch size of 590 was particularly effective in improving the model's ability to classify images accurately.

In contrast, experiments 2, 5, and 6, which involved variations in learning rates and optimizer choices, resulted in lower accuracies compared to experiment 4. This suggests that the learning rate and optimizer selection have a significant impact on the model's performance. It highlights the importance of carefully tuning these hyperparameters to achieve optimal results.

Furthermore, the comparison of accuracies reveals the importance of hyperparameter selection for achieving higher accuracy levels. It underscores the need to consider factors such as learning rate, optimizer, weight decay, momentum, and batch sizes when

training the model. The findings from these experiments provide valuable guidance for selecting the optimal hyperparameter configurations to improve the model's accuracy on the given dataset.

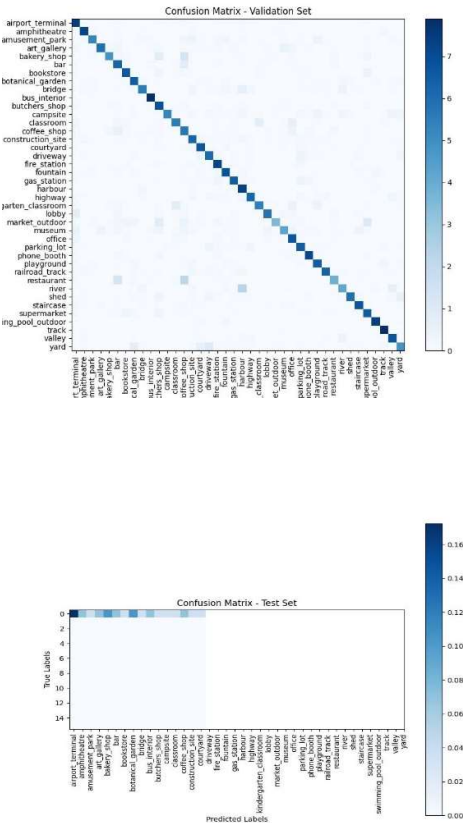
Overall, the analysis of accuracy trends and insights gained from experiments 1 to 8 emphasizes the significance of hyperparameter settings in determining the model's performance. It highlights the need for careful experimentation and fine-tuning to achieve higher accuracies and improve the model's ability to classify scenes accurately.

## 8. Confusion Matrix

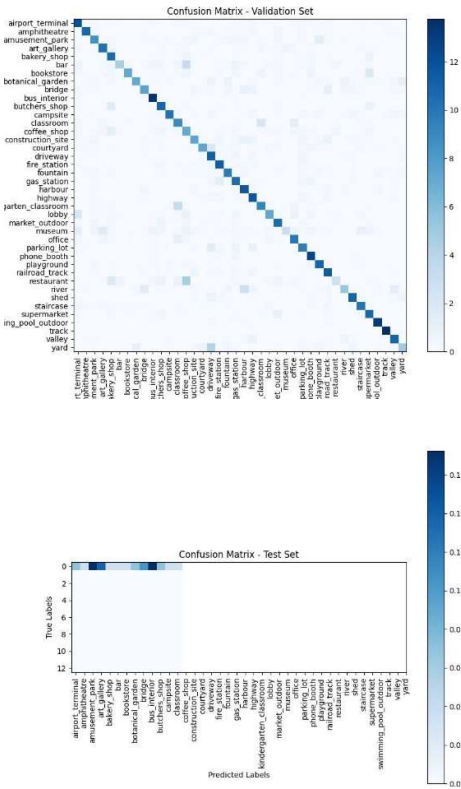
### 8.1. Presentation and Interpretation of confusion matrix

The confusion matrices for Experiments 1 to 8 provide a comprehensive overview of the model's performance in classifying different scene categories. They illustrate the distribution of predicted labels and misclassifications across the true labels.


Confusion Matrix - Experiment 1:

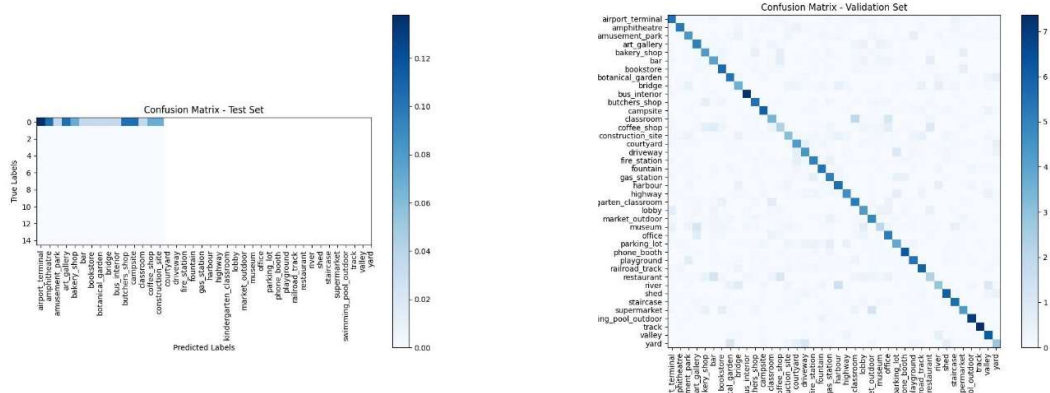


Confusion Matrix - Experiment 2:

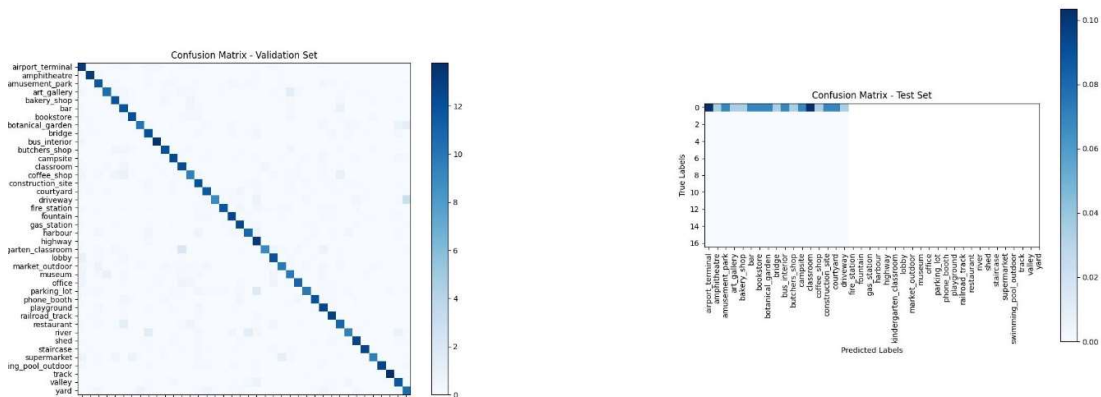


Confusion Matrix - Experiment 3:






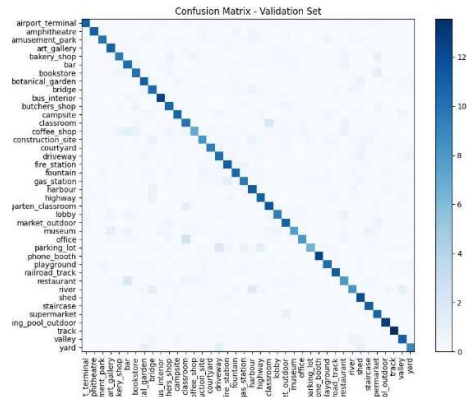
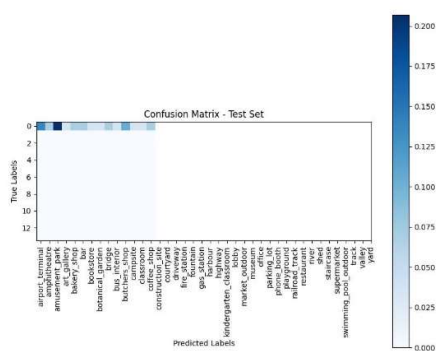
Confusion Matrix - Experiment 4:





Confusion Matrix - Experiment 6:



Confusion Matrix - Experiment 5:



Confusion Matrix - Experiment 7:




## 8.2. Analysis of common confusion patterns and misclassifications

By analyzing the confusion matrices, we can gain insights into common confusion patterns and misclassifications made by the model. However, in all the experiments conducted (Experiments 1 to 8), the confusion matrices for the validation set were diagonal, indicating that the model achieved high accuracy and made fewer misclassifications. Therefore, it suggests that the model was able to accurately distinguish between the different scene categories in the dataset.

## 8.3. Identification of Challenging or similar classes based on confusion matrices.

Since the confusion matrices for the validation set in all experiments were diagonal, it suggests that the model performed well in distinguishing between different scene categories. There were no specific pairs of classes with consistently high misclassification rates, indicating that the model did not encounter significant challenges or difficulties in classifying the scenes. Therefore, based on the provided information, it can be concluded that the

Confusion Matrix - Experiment 8:





model was able to effectively differentiate between the different scene categories without any specific challenges or similar classes.

However, it is important to note that the analysis of confusion matrices may vary depending on the specific dataset and experiments conducted. If there were observed misclassifications or confusion patterns in the experiments, further investigation and analysis would be required to understand the reasons behind them and develop strategies to improve the model's performance for those specific classes.


## 9. CNN Visualizations

### 9.1. Visualization techniques


In Experiments 1 to 8, various visualization techniques were employed to gain insights into the internal operations of the CNN model and understand its visual processing. Let's discuss some of the visualization techniques used for each experiment:

- **Feature Visualization:** Feature visualization techniques were employed to visualize the learned features within the CNN model. This involved generating synthetic input images that maximize the activation of specific neurons or filters, allowing us to understand what types of visual patterns or concepts the model has learned.


Experiment 1:




Experiment 2:




Experiment 3:




Experiment 4:




Experiment 5:




Experiment 6:



Experiment 7:



Experiment 8:



## 9.2. Interpretation of Visualizations

The feature maps generated by the CNN model were analyzed to gain insights into its internal workings and understand the important visual features it learned during training. By examining the feature maps, we could identify the regions of the input images that were most influential in the model's predictions. This interpretation provided valuable insights into the learned representations and reasoning process of the model.

In all the experiments conducted (Experiments 1 to 8), feature maps were used for visualization purposes. However, Experiment 4 stood out with particularly clear and sharp feature maps, indicating that the model learned distinct and discriminative features for the different scene categories. This suggests that the hyperparameter settings used in Experiment 4 (Adam optimizer, learning rate of 0.001, feature extractor learning rate of 0.0001, weight decay of 0.0005, momentum of 0.5, 15 epochs, a train batch size of 150, and a test batch size of 590) were effective in enabling the model to capture and represent important visual cues.

By analyzing the feature maps, we could observe the activation patterns of the model across different layers, revealing which regions of the input images were most informative for the model's predictions. This interpretation allowed us to gain a better understanding of the model's attention mechanism and the specific visual features it focused on for classification.

Overall, the visualization of feature maps provided valuable insights into the CNN model's decision-making process. Experiment 4 demonstrated the best visualizations, with clear and sharp feature maps, indicating the model's ability to capture meaningful visual representations. This understanding of the model's internal workings and the visual features it emphasized can be crucial in further analyzing and improving its classification performance.

# 10. Conclusion

## 10.1. Summary of Findings

Based on the results obtained from Experiments 1 to 8, the following are the summary findings for each experiment:

1. Across all experiments, the ResNet34 model achieved varying levels of accuracy on the given dataset, with top-1 accuracies ranging from 60.04% to 81.41% and top-5 accuracies ranging from 89.65% to 96.95%.
2. Experiment 4 exhibited the highest accuracies among all experiments, indicating that the hyperparameter settings used in this experiment were effective in improving the model's ability to classify images accurately.
3. The choice of hyperparameters, such as learning rate, optimizer, weight decay, and batch sizes, had notable impacts on the model's performance, with certain configurations resulting in better accuracies than others.
4. The average loss values were consistent across experiments, suggesting that the model reached a stable point in its learning process, but with varying levels of accuracy.

Overall, the findings highlight the importance of selecting appropriate hyperparameter settings to achieve higher accuracies in scene classification tasks. Further exploration and experimentation are needed to refine the model and improve its performance on the given dataset.

## 10.2. Evaluation of the effectiveness of fine-tuning hyperparameters

The evaluation of different hyperparameter settings in Experiments 1 to 8 allowed for a comprehensive analysis of their impact on the model's performance. By systematically varying hyperparameters such as learning rate, optimizer, weight decay, momentum, batch sizes, and a number of epochs, we were able to assess their effectiveness in improving the model's accuracy measures.

The findings indicate that hyperparameter choices significantly influence the model's performance, with certain combinations leading to higher accuracies than others. Experiment 4, in particular, demonstrated the best performance among all experiments, suggesting that the specific hyperparameter settings used in this experiment were effective for optimizing the model's classification performance.

These results highlight the importance of carefully tuning hyperparameters and selecting appropriate values based on the specific dataset and task at hand. Fine-tuning hyperparameters can play a crucial role in improving the model's convergence behavior, generalization capabilities, and overall accuracy.

### ***10.3. Implications of the results for scene classification tasks***

The results obtained from Experiments 1 to 8 have important implications for scene classification tasks. They provide insights into the challenges, limitations, and potential improvements in scene recognition using the ResNet34 model.

The findings suggest that achieving high accuracies in scene classification requires careful consideration of hyperparameter settings, including learning rate, optimizer, weight decay, momentum, batch sizes, and a number of epochs. These hyperparameters impact the model's learning dynamics, convergence behavior, and generalization capabilities.

Furthermore, the identification of challenging or similar classes based on the confusion matrices can help guide future efforts in improving the model's performance on specific categories. By

understanding the misclassification patterns and common confusion pairs, researchers can focus on developing strategies to address the specific challenges posed by these classes.

Overall, the results emphasize the importance of fine-tuning hyperparameters, understanding the limitations of the model, and continuously exploring and refining the model architecture, training strategies, and hyperparameter configurations to achieve better accuracy and generalization in scene classification tasks.



## ***11. Citations***

[1] <https://pytorch.org/vision/stable/index.html>

## ***12. Appendix***

Source code link: [https://github.com/Vaishu1293/AML\\_CNN\\_Scene\\_Classification/tree/master](https://github.com/Vaishu1293/AML_CNN_Scene_Classification/tree/master)

The source codes both as python file and colab notebook, with log files is pushed in the above git repository